

PS5_AndrewYu

July 29, 2023

```
[1]: from google.colab import drive

drive.mount('/content/drive')
FOLDERNAME = "Stanford Summer Session/SOC 128D"

#Packages for working with data
# !pip install "modin[dask]"
# import modin.pandas as pd
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
import plotnine as pn
import numpy as np

plt.style.use("ggplot")
```

Mounted at /content/drive

1. Run a topic model with 4 topics on the “guardian_co_river.csv” dataset.

```
[15]: import json
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk import word_tokenize
nltk.download('punkt')
from string import punctuation
from sklearn.feature_extraction.text import CountVectorizer
import regex as re
from nltk import pos_tag
nltk.download('averaged_perceptron_tagger')
from nltk.stem import WordNetLemmatizer
nltk.download('wordnet')

with open("stopwords_extra.json", "r") as file: stopwords_extra = set(json.
    ↪load(file))
stopwords_nltk = set(stopwords.words('english'))
punctuation += "-- ' , " " „ † ‡ • … % < > / "
```

```

stopwords_punct = set(punctuation)
combined_stoplist = list(set.union(stopwords_extra, stopwords_nltk,
    stopwords_punct))

# I used the dataset with 18548 articles, for practice working with bigger
    datasets
df = pd.read_csv(f"drive/My Drive/{FOLDERNAME}/data/guardian_co_river.csv").
    dropna(subset=["body_text"])
# Unicode Categories C (Control), M (Mark), P (Punctuation), S (Symbol), Z
    (Separator) + emojis
df["body_text"] = df["body_text"].apply(lambda x: re.
    compile(r'[\p{C}|\p{M}|\p{P}|\p{S}|\p{Z}]+', re.UNICODE).sub(" ", x).strip())

# Lemmatization
def get_POS_tags(pos_tag):
    POS_tag = {'NN': 'n', 'JJ': 'a', 'VB': 'v', 'RB': 'r'}
    try:
        # Getting first 2 letters of pos_tag
        return POS_tag[pos_tag[:2]]
    except:
        # Fallback to noun (Default)
        return 'n'

# In doing this, token_pattern parameter in CountVectorizer won't work, so just
    add any patterns here
class LemmaTokenizer(object):
    def __init__(self):
        self.wnl = WordNetLemmatizer()
    def __call__(self, docs):
        return [self.wnl.lemmatize(word, pos=get_POS_tags(tag)) for word, tag
    in pos_tag(word_tokenize(docs)) if not re.search("\d+", word)]

# This python library automatically integrates anti-joining stopwords and takes
    in corpus as list of texts, converting into list of words via word_tokenize,
    to output DTM. Thus unnest_tokens not used
tf_vectorizer = CountVectorizer(stop_words=combined_stoplist,
    tokenizer=LemmaTokenizer(), strip_accents=None)
# tf is in the form (doc, word): freq
tf = tf_vectorizer.fit_transform(df["body_text"])

print(f"No. of words per topic: {len(tf_vectorizer.get_feature_names_out())}")
pd.DataFrame({"word": tf_vectorizer.get_feature_names_out()})

```

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!

```



```

else:
    for topic_i, topic_freq in enumerate(model.components_):
        top_freq_words_i = topic_freq.argsort()[::-n_top_words - 1:-1]
        # {topic: [word, ...]}
        word_dict[topics[topic_i]] = [words[i] for i in top_freq_words_i]
    return pd.DataFrame(word_dict)

from sklearn.decomposition import LatentDirichletAllocation as LDA
lda4 = LDA(n_components=4, random_state=1)
# probability of each (4) topic per (17982) document
topic_per_document = lda4.fit_transform(tf)
# Frequency of each (157463) word-term per (4) topic
word_per_topic = lda4.components_

# Log Likelihood: Less negative
print("Log Likelihood: ", lda4.score(tf))

# Perplexity: exp(-1 * log likelihood), Lower
print("Perplexity: ", lda4.perplexity(tf))

topics_words_df4 = get_model_topics(lda4, tf_vectorizer, ["Topic 1", "Topic 2",
↳ "Topic 3", "Topic 4"])
topics_words_df4

```

Log Likelihood: -86926244.37087119

Perplexity: 5463.206371294159

```

[17]:
      Topic 1      Topic 2 Topic 3 Topic 4
0      people      trump   year   year
1      water      state   make   day
2      year republican people  city
3 government president   time  river
4      state      win    work   time
5      report      vote   life  world
6      country      make  river  park
7      river      year   find   make
8      case      time  child   run
9      make      election film   back

```

2. Plot the top words for each topic, similar to what is done in section 5.2. What would you label these topics?

```

[104]: topics_words_df4 = get_model_topics(lda4, tf_vectorizer, ["Topic 1", "Topic 2",
↳ "Topic 3", "Topic 4"], detailed=True)
topics_words_df4

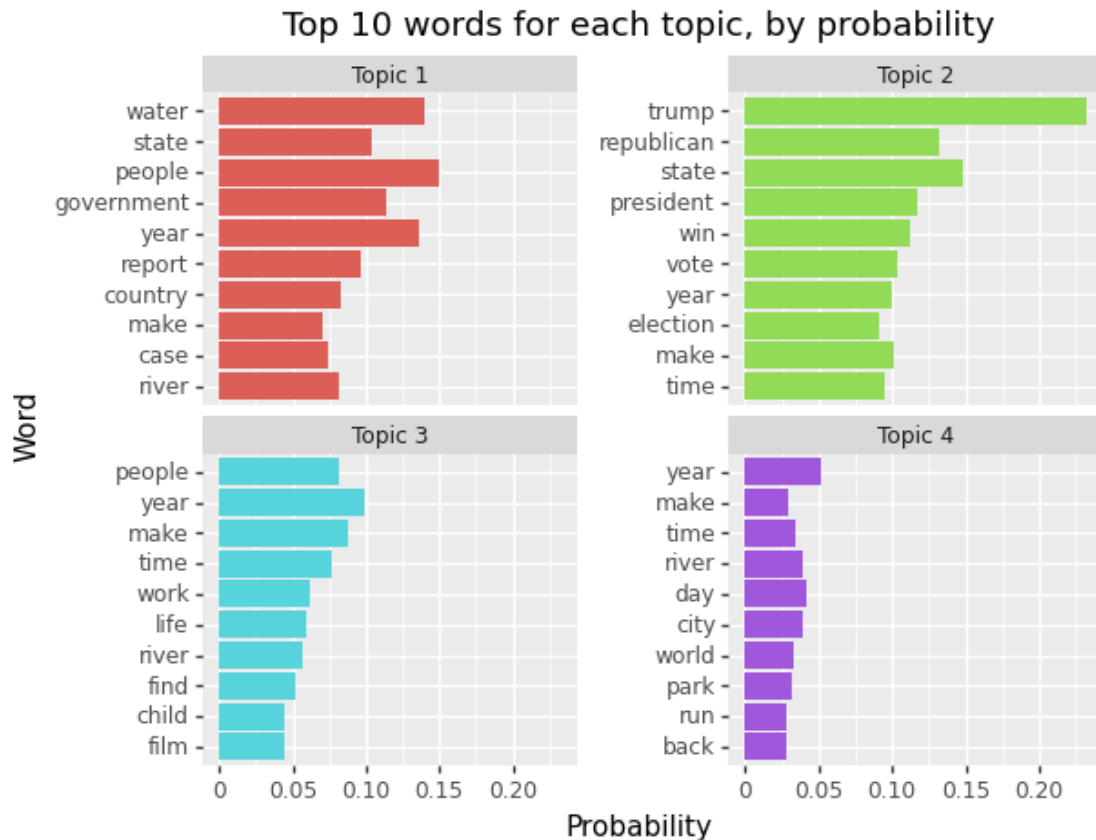
```

```
[104]:
```

	Topic	Word	Probability
0	Topic 1	people	0.150188
1	Topic 1	water	0.139437
2	Topic 1	year	0.135564
3	Topic 1	government	0.113529
4	Topic 1	state	0.104199
5	Topic 1	report	0.096412
6	Topic 1	country	0.082658
7	Topic 1	river	0.081916
8	Topic 1	case	0.074222
9	Topic 1	make	0.070491
10	Topic 2	trump	0.232615
11	Topic 2	state	0.147654
12	Topic 2	republican	0.132483
13	Topic 2	president	0.117492
14	Topic 2	win	0.112478
15	Topic 2	vote	0.102859
16	Topic 2	make	0.101298
17	Topic 2	year	0.099786
18	Topic 2	time	0.095326
19	Topic 2	election	0.091645
20	Topic 3	year	0.098372
21	Topic 3	make	0.087154
22	Topic 3	people	0.081191
23	Topic 3	time	0.076403
24	Topic 3	work	0.061427
25	Topic 3	life	0.059303
26	Topic 3	river	0.057298
27	Topic 3	find	0.051575
28	Topic 3	child	0.044530
29	Topic 3	film	0.044246
30	Topic 4	year	0.051727
31	Topic 4	day	0.041634
32	Topic 4	city	0.039051
33	Topic 4	river	0.038604
34	Topic 4	time	0.033675
35	Topic 4	world	0.033473
36	Topic 4	park	0.031982
37	Topic 4	make	0.028890
38	Topic 4	run	0.028120
39	Topic 4	back	0.027374

```
[105]: # I've tried searching but I'm not sure why the sorting is not working, df is
        ↪sorted, and x is also reordered within
plot = (pn.ggplot(topics_words_df4, pn.aes(x='reorder(Word, Probability)',
        ↪y='Probability', fill='Topic')) +
        pn.geom_col(show_legend=False) +
```

```
pn.coord_flip() +
pn.facet_wrap('Topic', scales = "free_y") +
pn.theme(subplots_adjust={'wspace': 0.4}) +
pn.labs(title="Top 10 words for each topic, by probability", x="Word", y="Probability")
print(plot)
```



They are words are not very distinct, due to choice of number of topics, but I would label them

Topic 1: Government

Topic 2: Politics

Topic 3: Work

Topic 4: Time

3. Try running the LDA again, this time with a different number of topics. Does the new topic model make more or less sense to you?

```
[ ]: from sklearn.model_selection import GridSearchCV

lda = LDA()
model = GridSearchCV(lda, param_grid={'n_components': [2, 4, 6, 8, 10]})
model.fit(tf)

print("Best Model's Params: ", model.best_params_)
print("Best Log Likelihood Score: ", model.best_score_)
print("Model Perplexity: ", model.best_estimator_.perplexity(tf))
```

```
Best Model's Params: {'n_components': 2}
Best Log Likelihood Score: -20246738.786500275
Model Perplexity: 9943.179337936572
```

```
[19]: lda2 = LDA(n_components=2, random_state=1)
_ = lda2.fit_transform(tf)

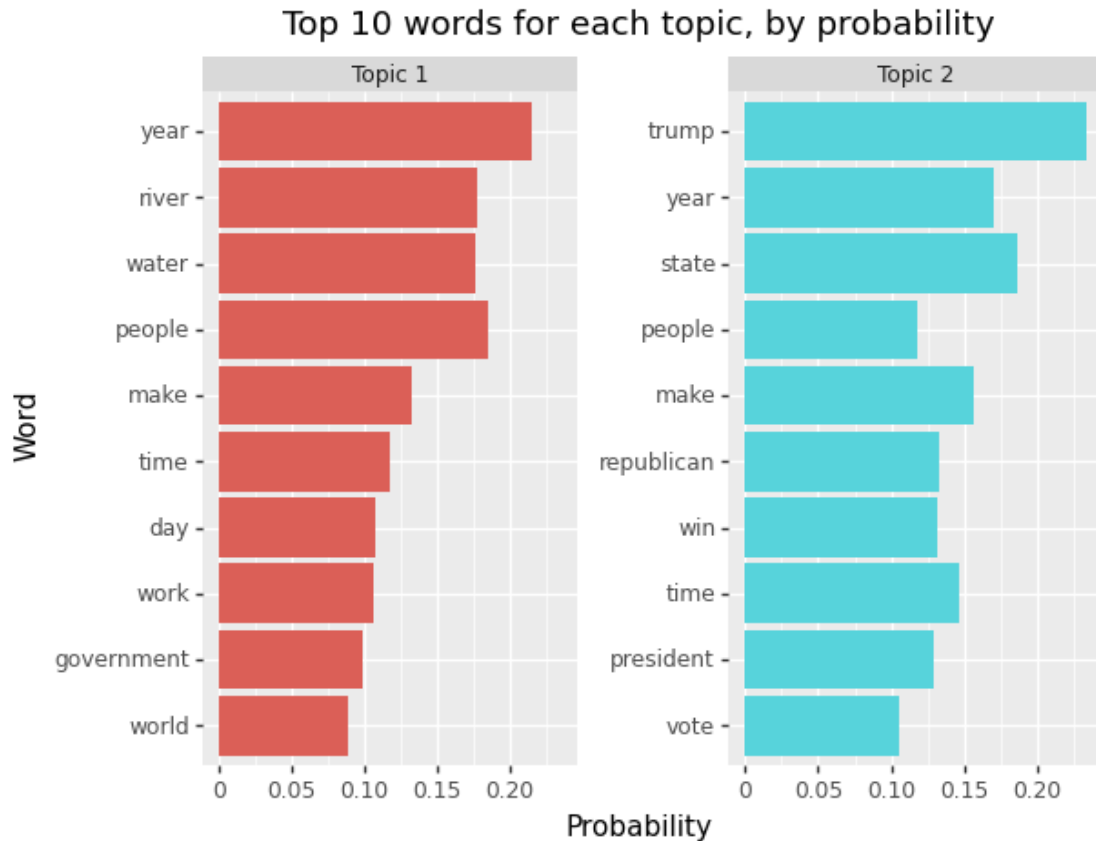
topics_words_df2 = get_model_topics(lda2, tf_vectorizer, ["Topic 1", "Topic 2"])
topics_words_df2
```

```
[19]:
```

	Topic 1	Topic 2
0	year	trump
1	people	state
2	river	year
3	water	make
4	make	time
5	time	republican
6	day	win
7	work	president
8	government	people
9	world	vote

```
[107]: topics_words_df2 = get_model_topics(lda2, tf_vectorizer, ["Topic 1", "Topic_
↳2"], detailed=True)
plot = (pn.ggplot(topics_words_df2, pn.aes(x='reorder(Word, Probability)',
↳y='Probability', fill='Topic')) +
        pn.geom_col(show_legend=False) +
        pn.coord_flip() +
        pn.facet_wrap('Topic', scales = "free_y") +
        pn.theme(subplots_adjust={'wspace': 0.4}) +
        pn.labs(title="Top 10 words for each topic, by probability", x="Word",
↳y="Probability")
        )

print(plot)
```



My intuition was 2 topics (based on the in-class materials), but I also wanted to test the optimal number using the GridSearch, which confirmed my intuition.

As evident, 2 topics summarises all the documents more clearly and sensibly, with Topic 1 being about rivers and Topic 2 being more political

- Using whichever topic model made the most sense to you, create a dataset similar to `co_documents` in section 5.3. Choose two articles to examine in more detail. Discuss the topic probabilities.

```
[123]: def get_model_docs(model, dtm, detailed=False):
    topic_dict = {}
    doc_topic = model.transform(dtm)
    if detailed:
        for doc_i, topic_probs_of_doc in enumerate(doc_topic):
            num_topics = doc_topic.shape[1]
            # Sorting topics by top probabilities
            top_probable_topics = topic_probs_of_doc.argsort()[::-num_topics - 1:-1]
            topic_dict[f"Doc {doc_i}"] = [(topic+1, topic_probs_of_doc[topic]) for
            ↪topic in top_probable_topics]
```



```

df = pd.DataFrame([(doc, topic, prob) for doc, topics_probs in topic_dict.
↳items() for topic, prob in topics_probs], columns=["Document", "Topic",
↳"Probability"])

# Chaining doesn't work for some reason
df['Dominant_topic'] = df[df.index % num_topics == 0]["Topic"]
df['Dominant_topic'] = df['Dominant_topic'].fillna(method='ffill')
df['Dominant_topic'] = df['Dominant_topic'].astype(int)
return df

else:
    for doc_i, topic_probs_of_doc in enumerate(doc_topic):
        num_topics = doc_topic.shape[1]
        top_probable_topics = topic_probs_of_doc.argsort()[:-num_topics - 1:-1]
        topic_dict[f"Doc {doc_i}"] = [topic for topic in top_probable_topics]
    return pd.DataFrame(topic_dict)

docs_topics_df = get_model_docs(lda2, tf, detailed=True)
docs_topics_df
# Probabilities of each topic add up to 1, which confirms this is similar to
↳co_documents (gamma)

```

```

[123]:
      Document  Topic  Probability  Dominant_topic
0         Doc 0      1      0.645597             1
1         Doc 0      2      0.354403             1
2         Doc 1      1      0.999263             1
3         Doc 1      2      0.000737             1
4         Doc 2      1      0.921166             1
...         ...    ...          ...             ...
35959  Doc 17979      2      0.004749             1
35960  Doc 17980      2      0.987856             2
35961  Doc 17980      1      0.012144             2
35962  Doc 17981      1      0.952474             1
35963  Doc 17981      2      0.047526             1

```

[35964 rows x 4 columns]

```

[125]: # To find a topic that has a very fair probability of belonging to either Topic
docs_topics_df[docs_topics_df["Probability"].between(0.45, 0.65)]
# I chose Doc 17975

```

```

[125]:
      Document  Topic  Probability  Dominant_topic
0         Doc 0      1      0.645597             1
18        Doc 9      1      0.593177             1
48        Doc 24     1      0.539227             1
49        Doc 24     2      0.460773             1
56        Doc 28     2      0.571738             2
...         ...    ...          ...             ...

```

35902	Doc 17951	2	0.586099	2
35938	Doc 17969	2	0.555969	2
35946	Doc 17973	1	0.635390	1
35950	Doc 17975	2	0.506625	2
35951	Doc 17975	1	0.493375	2

[2900 rows x 4 columns]

```
[124]: chosen_docs_df = docs_topics_df[docs_topics_df["Document"].isin(["Doc 1", "Doc_
↪17975", "Doc 17980"])]
chosen_docs_df
```

```
[124]:
```

	Document	Topic	Probability	Dominant_topic
2	Doc 1	1	0.999263	1
3	Doc 1	2	0.000737	1
35950	Doc 17975	2	0.506625	2
35951	Doc 17975	1	0.493375	2
35960	Doc 17980	2	0.987856	2
35961	Doc 17980	1	0.012144	2

The probabilities show that

- Doc 1 is extremely likely to be Topic 1 (Rivers)
 - Doc 17975 is quite on the fence, but ultimately leans towards Topic 2 (Politics)
 - Doc 17980 is extremely likely to be Topic 2 (Politics)
5. Considering what we know about the Colorado River and water rights in the U.S. (from readings and class), are we surprised by the topics generated by these news articles? Describe how your results compared to expectations.

```
[131]: for index, row in df.iloc[[1, 17975, 17980], df.columns.
↪get_indexer(['web_title', 'web_url'])].iterrows():
print(f"Title: {row['web_title']}, URL:{row['web_url']}\n")
```

Title: The farmers dealing with water shortages even before historic Colorado River deal, URL:<https://www.theguardian.com/global/2023/may/31/arizona-farmers-water-colorado-river-cuts>

Title: Paraguayan dictator toppled in family coup,
URL:<https://www.theguardian.com/theguardian/1989/feb/04/fromthearchive>

Title: The American Presidency,
URL:<https://www.theguardian.com/world/1912/nov/07/usa.fromthearchive>

Upon reading the articles, the most probable topics mostly makes sense, as also suggested by the title of the article.

Doc 1 discusses the lives of the farmers near the Colorado River, which checks out with it being assigned Topic 1 (Rivers)

Doc 17975 discusses the the events surrounding the political party in Colorado, and has little to do with the Colorado River, yet is assigned such a fair probability, which is confusing. This could be because even Topic 1 itself isn't such a clearly defined "Rivers" topic - there are also quite a lot of political terms. What is the true differentiating factor between Topic 1 and 2 is that 2 is more specific ("trump", "republican", "president"). However, the model did predict correctly that this article is more likely Topic 2 than 1 anyway.

Doc 17980 discusses mostly the American Presidency which checks out with it being assigned Topic 2 (Politics)

Doc 1, which talks the most regarding the Colorado River situation, aligns with what we learnt from the readings - Arizona is fighting for its rights to usage of the water from the river, with California and Nevada, aligning with the [reading](#) - The farmers were already struggling with water, in growing the most profitable and marketable commodity crops - alfalfa and cotton. This aligns with the [reading](#) that agriculture sucks up 79% of all water usage from the river.

```
[ ]: # If it contains spaces, use '\ ' to represent each space E.g. 'Summer\ PSet\ 1.
      ↳ipynb'
FILENAME = "PS5_AndrewYu.ipynb"

%cd drive/My\ Drive
%cd $FOLDERNAME
!sudo apt-get install texlive-xetex texlive-fonts-recommended
      ↳texlive-plain-generic
!pip install PyPDF2

!jupyter nbconvert --log-level CRITICAL --to pdf $FILENAME
```

```
/content/drive/My Drive
/content/drive/My Drive/Stanford Summer Session/SOC 128D
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
```

The following additional packages will be installed:

```
dvisvgm fonts-droid-fallback fonts-lato fonts-lmodern fonts-noto-mono
fonts-texgyre fonts-urw-base35 libapache-pom-java libcommons-logging-java
libcommons-parent-java libfontbox-java libfontenc1 libgs9 libgs9-common
libidn12 libijs-0.35 libjbig2dec0 libkpathsea6 libpdfbox-java libptexenc1
libruby3.0 libsynchronet2 libteckit0 libtexlua53 libtexluajit2 libwoff1
libzip-0-13 lmodern poppler-data preview-latex-style rake ruby
ruby-net-telnet ruby-rubygems ruby-webrick ruby-xmlrpc ruby3.0
rubygems-integration tlutils teckit tex-common tex-gyre texlive-base
texlive-binaries texlive-latex-base texlive-latex-extra
texlive-latex-recommended texlive-pictures tipa xfonts-encodings
xfonts-utils
```

Suggested packages:

```
fonts-noto fonts-freefont-otf | fonts-freefont-ttf libavalon-framework-java
libcommons-logging-java-doc libexcalibur-logkit-java liblog4j1.2-java
poppler-utils ghostscript fonts-japanese-mincho | fonts-ipafont-mincho
```